

Tutorial 6 - Neural Network

Part B: classification on real world dataset

Oct, 2023

Yuan Zhong



Official website: <https://pytorch.org/get-started/locally/>

Importing The Required Libraries

```
In [5]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os

import torch
import torch.nn as nn
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [6]: df1 = pd.read_csv('./data/Iris.csv')
```

```
In [7]: df1.head()
```

```
Out[7]:
```

| | Id | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | Species |
|---|----|---------------|--------------|---------------|--------------|-------------|
| 0 | 1 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 2 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 3 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |

Creating A Function To Get The Details of The Dataset

```
In [8]: def get_info_dataframe(dataframe):
        print(f"DATAFRAME GENERAL INFO - \n")
        print(dataframe.info(), "\n")
        print(f"DATAFRAME MISSING INFO - \n")
        print(dataframe.isnull().sum(), "\n")
        print(f"DATAFRAME SHAPE INFO - \n")
        print(dataframe.shape)
```

```
get_info_dataframe(df1)
```

DATAFRAME GENERAL INFO -

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Id              150 non-null   int64
 1   SepalLengthCm   150 non-null   float64
 2   SepalWidthCm    150 non-null   float64
 3   PetalLengthCm   150 non-null   float64
 4   PetalWidthCm    150 non-null   float64
 5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None
```

DATAFRAME MISSING INFO -

```
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
PetalWidthCm    0
Species         0
dtype: int64
```

DATAFRAME SHAPE INFO -

```
(150, 6)
```

```
In [9]: df1['Species'].unique()
```

```
Out[9]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

LabelEncoding The Attributes of The Target Column

```
In [10]: df1['Species'] = df1['Species'].map({'Iris-setosa':0, 'Iris-versicolor':1, 'Iris-virginica':2})
```

```
In [11]: df1.head()
```

```
Out [11]:
```

| | Id | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | Species |
|---|----|---------------|--------------|---------------|--------------|---------|
| 0 | 1 | 5.1 | 3.5 | 1.4 | 0.2 | 0 |
| 1 | 2 | 4.9 | 3.0 | 1.4 | 0.2 | 0 |
| 2 | 3 | 4.7 | 3.2 | 1.3 | 0.2 | 0 |
| 3 | 4 | 4.6 | 3.1 | 1.5 | 0.2 | 0 |
| 4 | 5 | 5.0 | 3.6 | 1.4 | 0.2 | 0 |

```
In [12]: df1.drop(['Id'],axis=1,inplace=True)
```

```
In [13]: df1.head()
```

```
Out [13]:
```

| | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | Species |
|---|---------------|--------------|---------------|--------------|---------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | 0 |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | 0 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | 0 |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | 0 |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | 0 |

```
In [14]: X = df1.drop(["Species"],axis=1).values
y = df1["Species"].values
```

```
In [15]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
In [16]: scaler = StandardScaler()
```

Doing The Train Test Split And Scaling The Data

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, ra
```

```
In [33]: X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
print(X_train.shape)
```

```
(105, 4)
```

Converting From Numpy Array To Torch Tensor

```
In [19]: X_train = torch.FloatTensor(X_train)
```

```
X_test = torch.FloatTensor(X_test)
y_train = torch.LongTensor(y_train)
y_test = torch.LongTensor(y_test)
```

Creating Our Neural Network Model For Classification

```
In [20]: class NeuralNetworkClassificationModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(NeuralNetworkClassificationModel, self).__init__()
        self.input_layer = nn.Linear(input_dim, 128)
        self.hidden_layer1 = nn.Linear(128, 64)
        self.output_layer = nn.Linear(64, output_dim)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.relu(self.input_layer(x))
        out = self.relu(self.hidden_layer1(out))
        out = self.output_layer(out)
        return out
```

```
In [21]: # input_dim = 4 because we have 4 inputs namely sepal_length, sepal_width, petal_length, and petal_width
# output_dim = 3 because we have 3 categories setosa, versicolor, and virginica
input_dim = 4
output_dim = 3
model = NeuralNetworkClassificationModel(input_dim, output_dim)
```

```
In [22]: # creating our optimizer and loss function object
learning_rate = 0.01
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

```
In [24]: def train_network(model, optimizer, criterion, X_train, y_train, X_test, y_test, num_epochs):
    for epoch in range(num_epochs):
        #clear out the gradients from the last step loss.backward()
        optimizer.zero_grad()

        #forward feed
        output_train = model(X_train)

        #calculate the loss
        loss_train = criterion(output_train, y_train)

        #backward propagation: calculate gradients
        loss_train.backward()

        #update the weights
        optimizer.step()

    output_test = model(X_test)
```

```

        loss_test = criterion(output_test,y_test)

        train_losses[epoch] = loss_train.item()
        test_losses[epoch] = loss_test.item()

        if (epoch + 1) % 50 == 0:
            print(f"Epoch {epoch+1}/{num_epochs}, Train Loss: {loss_train.it

```

```

In [25]: num_epochs = 1000
         train_losses = np.zeros(num_epochs)
         test_losses = np.zeros(num_epochs)

```

```

In [26]: train_network(model,optimizer,criterion,X_train,y_train,X_test,y_test,num_ep

```

```


Epoch 50/1000, Train Loss: 0.0318, Test Loss: 0.0034
Epoch 100/1000, Train Loss: 0.0018, Test Loss: 0.0875
Epoch 150/1000, Train Loss: 0.0004, Test Loss: 0.0757
Epoch 200/1000, Train Loss: 0.0002, Test Loss: 0.0818
Epoch 250/1000, Train Loss: 0.0001, Test Loss: 0.0851
Epoch 300/1000, Train Loss: 0.0001, Test Loss: 0.0878
Epoch 350/1000, Train Loss: 0.0001, Test Loss: 0.0877
Epoch 400/1000, Train Loss: 0.0001, Test Loss: 0.0894
Epoch 450/1000, Train Loss: 0.0000, Test Loss: 0.0894
Epoch 500/1000, Train Loss: 0.0000, Test Loss: 0.0899
Epoch 550/1000, Train Loss: 0.0000, Test Loss: 0.0902
Epoch 600/1000, Train Loss: 0.0000, Test Loss: 0.0905
Epoch 650/1000, Train Loss: 0.0000, Test Loss: 0.0908
Epoch 700/1000, Train Loss: 0.0000, Test Loss: 0.0911
Epoch 750/1000, Train Loss: 0.0000, Test Loss: 0.0911
Epoch 800/1000, Train Loss: 0.0000, Test Loss: 0.0909
Epoch 850/1000, Train Loss: 0.0000, Test Loss: 0.0910
Epoch 900/1000, Train Loss: 0.0000, Test Loss: 0.0909
Epoch 950/1000, Train Loss: 0.0000, Test Loss: 0.0909
Epoch 1000/1000, Train Loss: 0.0000, Test Loss: 0.0908

```

```

In [27]: plt.figure(figsize=(10,10))
         plt.plot(train_losses, label='train loss')
         plt.plot(test_losses, label='test loss')
         plt.legend()
         plt.show()

```



```
In [28]: predictions_train = []
         predictions_test = []
         with torch.no_grad():
             predictions_train = model(X_train)
             predictions_test = model(X_test)
```

```
In [29]: # Check how the predicted outputs look like and after taking argmax compare
         #predictions_train
         #y_train,y_test
```

```
In [30]: def get_accuracy_multiclass(pred_arr,original_arr):
         if len(pred_arr)!=len(original_arr):
             return False
         pred_arr = pred_arr.numpy()
         original_arr = original_arr.numpy()
         final_pred= []
         # we will get something like this in the pred_arr [32.1680,12.9350,-58.4
         # so will be taking the index of that argument which has the highest val
         for i in range(len(pred_arr)):
             final_pred.append(np.argmax(pred_arr[i]))
```

```
final_pred = np.array(final_pred)
count = 0
#here we are doing a simple comparison between the predicted_arr and the
for i in range(len(original_arr)):
    if final_pred[i] == original_arr[i]:
        count+=1
return count/len(final_pred)
```

```
In [31]: train_acc = get_accuracy_multiclass(predictions_train,y_train)
test_acc = get_accuracy_multiclass(predictions_test,y_test)
```

```
In [32]: print(f"Training Accuracy: {round(train_acc*100,3)}")
print(f"Test Accuracy: {round(test_acc*100,3)}")
```

```
Training Accuracy: 100.0
Test Accuracy: 97.778
```